

Robust and Generalizable Decision-Making for CAVs at Unsignalized Intersections

A Deep Reinforcement Learning Approach Reproduction, Benchmarking & Path to
Generalization

Mahmudul Hasan (220041231)

Ishmam Tahmid (220041259)

Farhan Tahsin Khan (220041229)

CSE 4610: Design Project

Supervisors: Ashraful Alam Khan & Tanjila Alam Sathi


Course Coordinator: Dr. Md. Azam Hossain

September 14, 2026

Outline

- 1 Introduction
- 2 Related Work
- 3 Methodology
- 4 Implementation
- 5 Experimental Results
- 6 Challenges and Limitations
- 7 Conclusion and Future Direction
- 8 Project Repository

The Problem




36%

of crashes are
intersection-related



\$371B

annual societal cost
of intersection crashes




40%+

of collisions occur
at intersections in
International Context

- Traffic lights control at the [lane-group](#) level, not per vehicle
- Fixed timing → stopping even when the road is clear
- Unsignalized crossings: high risk from occlusion & uncertain intent
- No shared protocol → right-of-way must be inferred

Motivation: AV vs. CAV



Signal-Free AIM

Individual, continuous negotiation replaces group-level stopping

Deep RL

Learns adaptive go/yield behavior under stochastic traffic

How should an ego vehicle decide, every timestep, to **go** or **yield** among randomly-behaving vehicles — safely, efficiently, reliably?

- Simultaneous, uncertain decisions among interacting vehicles
- No reservation protocol — right-of-way is inferred, not commanded
- Must generalize across density, driving style, and (eventually) geometry

Key Objectives

Safety

Zero (or near-zero) collisions
across randomized turning
tasks


Efficiency

Minimize crossing / travel
time, maximize throughput
speed

Robustness

Stable performance across
traffic densities *and*,
eventually, geometries

Taxonomy of Coordination Strategies



What the Literature Shows

Axis	Most popular	Runner-up
Goal	Safety + Efficiency (37%)	Safety alone (26%)
Architecture	V2V (159 papers)	V2I (119 papers)
Scheduling	Optimization (277 papers)	Heuristic (20 papers)
Modeling	Trajectory Planning	Spatio-Temporal Reservation

RL Literature Thread I: Single-Ego Decision Making

Paper	Idea	Architecture	Key finding
Isele et al. (2018)	DQN vs. TTC + occlusion “creeping”	Single-agent, no comm.	28% faster, <1% collisions
Wu et al. (2019)	Decentralized coordination Q-learning	Multi-agent, sparse comm.	Beats FCFS/LQF/Webster
Kai et al. (2020)	One policy, 3 maneuvers via task vector	Single-agent	Matches/beats 3 separate DQNs
Shu et al. (2022)	Weight transfer across maneuvers	Single-agent	Same-task 100%, cross-task >70%


RL Literature Thread II: Graph Representations & Cooperative AIM

Paper	Contribution	Setting
Chen et al. (2025)	Scenario-agnostic “lane-gap” + PPO + CILQR	Single-ego, multi-scenario (93–97%)
Schier et al. (2023)	Vehicle+road graph, learned edges	Predicts right-of-way at 99.5% acc.
Klimke et al. (2022)	Edge-augmented RGCN + TD3, controls <i>all</i> vehicles	True cooperative AIM
Huamán Reyna et al. (2026)	Survey, 53 papers	Hybrid methods score best overall

Where we fit

- Single-ego decision-making, multi-task state (like Kai et al., Xiao et al.)
- Not centralized AIM (Klimke et al.) — tractable scope, still comparable

Methodology at a Glance



Each stage builds on the previous milestone's validated pipeline — detailed next.

Reproduction Target: Xiao et al. (2024, IEEE T-VT)

“Decision-Making for AVs in Random Task Scenarios at Unsignalized Intersection Using Deep RL”

Three contributions

1. **Mix-Attention Network** (replaces plain MLP)
2. **Task-intent variable** μ_a
3. **Segmented replay buffer**

Setting

4-way, single-lane, unsignalized; random task/episode; **highway-env**; Discrete SAC, 50,000 episodes

Disclaimer: Scope of This Work

Reproduction & benchmarking study — not our primary contribution

Build a working pipeline, understand the environment deeply, then propose a novel direction.

Acknowledged limitations

- Original source code **not released**
- Some hyperparameters / reward constants only partially specified
- Reproduction is faithful, but **not byte-for-byte**

Research Progression: Three Milestones

M1 — Basic Ring Stabilization

Single-agent PPO, 230m ring, 21 IDM cars

Result: damped phantom shockwaves



M2 — Phase 1 Exploratory Reproduction

PPO, 4-way intersection, Original vs. Config B ($+\mu_a$)

Result: found left-turn vulnerability; fixed eval bug



M3 — Phase 2 Exact Reproduction

Discrete SAC (twin critics, auto-tuned α), 50k episodes

Result: 97.8% success, 0% collisions (1,050 tests)

State Space Formulation

Observation vector (48 m radar radius, up to 9 nearest surrounding vehicles):

$$\mathbf{S} = [S_{\text{ego}}, S_1, S_2, \dots, S_9]$$

Ego Vehicle State ($S_{\text{ego}} \in \mathbb{R}^{6 \text{ or } 7}$)

$$\mathbf{S}_{\text{ego}} = [p_{\text{ego}}, x, y, v_x, v_y, \phi, (\mu_a)]$$

- $p_{\text{ego}} \in \{0, 1\}$: **Existence flag** (1 = present)
- x, y : **Global position** coordinates (m)
- v_x, v_y : **Absolute velocity** along x, y (m/s)
- ϕ : **Heading angle** / orientation (rad)
- μ_a : **Task-intent feature** (*Config B only*)
Encodes route (left/straight/right) & exit progress

9 Surrounding Vehicles ($S_i \in \mathbb{R}^6$ each)

$$\mathbf{S}_i = [p_i, \Delta x_i, \Delta y_i, \Delta v_{x,i}, \Delta v_{y,i}, \phi_i]$$

- $p_i \in \{0, 1\}$: **Presence flag** (0 if empty / padded)
- $\Delta x_i, \Delta y_i$: **Relative position** ($x_i - x, y_i - y$) (m)
- $\Delta v_{x,i}, \Delta v_{y,i}$: **Relative velocity**
($v_{x,i} - v_x, v_{y,i} - v_y$) (m/s)
- ϕ_i : **Heading angle** of vehicle i (rad)
- **Ordering**: Sorted ascending by distance
 $d_i \leq 48$ m

Action Space & Reward Function

Action (discrete, 3 values)

$$A = \{\text{accel}, \text{idle}, \text{decel}\}$$

$$a = K_{P,A}(v_{tar} - v), \quad K_{P,A} = 5$$


Lateral control is a fixed PD controller
— RL controls only longitudinal speed.

Reward

$$R = \begin{cases} +30 & \text{reach destination} \\ -120 & \text{collision} \\ +0.1 & \text{reach target speed} \end{cases}$$

- Sparse terminal terms (safety + completion)
- Small dense term rewards cruising
- Collision penalty $\gg$ arrival reward

System Pipeline



Discrete Soft Actor-Critic (Discrete SAC)


Why SAC over PPO

- Off-policy $\rightarrow$ sample efficient
- Max-entropy $\rightarrow$ broad exploration
- Auto-tuned α — no fragile hyperparameter

$$\pi^* = \arg \max_{\pi} \sum_n \mathbb{E}[\gamma^n (R_n + \alpha \mathcal{H}(\pi(\cdot|S_n)))]$$

- Twin Q-critics reduce overestimation bias
- Discrete policy: softmax over 3 actions


Implementation Stack



Simulation & Training Setup


	Project 1: Ring	Project 2: Phase 1	Project 3: Phase 2	Paper
Scenario	230m ring	4-way	4-way	4-way
Sim. engine	Flow + SUMO	Flow + SUMO	highway-env	highway-env
Algorithm	PPO	PPO	Discrete SAC	Discrete SAC
Network	Actor-Critic MLP	[32, 32, 32]	[256, 256, 64, 3]	[256, 256, 64, 3]
Obs. space	3-dim POMDP	60 / 61-dim	60 / 61-dim	60 / 61-dim
Action space	Continuous [−1, 1] m/s ²	Discrete(3)	Discrete(3)	Discrete(3)
Training budget	123 iters	2.5k / 9.36k iters	50k episodes	50k episodes

Result 1: Success & Collision Rate




97.8% success, 0% collisions over 1,050 episodes — exceeds the paper's 97.4%/2.6%.

Result 2: Mean Crossing Speed



Nearly **2× faster** crossing than the paper — the simulators are not physically equivalent.

Result 3: Intent-Driven Speed Modulation



Config B ($+\mu_a$) slows for turns (Left 6.91, Right 6.32 m/s); Original stays uniform (~ 7.00 m/s).

Challenges and Limitations

Incomplete Specification

Original paper code & constants unreleased; required extensive parameter inference.

Multi-Ego Evaluation Bug

Caught and resolved an evaluation reward attribution bug during Milestone 2.

Partial Coverage

Validated 2 of 8 paper configurations; remaining 6 ready for future extension.

Hardware Bottleneck

Hybrid-CPU thread-sync stalls resolved via physical P-core affinity pinning ($3.5\times$ speedup).

Sim-to-Sim Physics Gap

Simulator collision kinematics diverge between `highway-env` and paper platform.

Single Geometry

Validated solely on standard 4-way intersection; motivation for geometry-agnostic model.

Proposed Novelty: Geometry-Agnostic Decision-Making


Core direction

Every model above is trained/evaluated on *one* fixed 4-way geometry — real deployments include T-junctions, merges, roundabouts.

Goal: a policy that does not depend heavily on its training geometry.


- Evaluate: 4-way, T-junction, ramp merge, roundabout
- Measure transfer of right-of-way reasoning, not just re-training

Generalization Strategy



Graph-based and relational representations generalize far better across unseen junction layouts than fixed-coordinate vectors.



Proposed Project Timeline



Conclusion

- Full DRL pipeline validated across three progressive milestones
- 1:1 Discrete SAC reproduction, matching and exceeding published benchmarks
- Performance gap traced to simulator collision physics and specification gaming
- Concrete path and experimental framework established toward [geometry-agnostic](#) decision-making

GitHub Contributors Activity Graph



Thank You!

Questions & Discussion

Project Repository

<https://github.com/Farhan41229/Drl-unsignalized-intersection-navigation>

References I



Liang Chen et al.

Unified planning framework with drivable area attention extraction for autonomous driving in urban scenarios.
IEEE Robotics and Automation Letters (RA-L), 2025.



Ashkan Gholamhosseinian and Jochen Seitz.

A comprehensive survey on cooperative intersection management for heterogeneous connected vehicles.
IEEE Access, 10:7937–7972, 2022.



Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine.

Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor.
In *International Conference on Machine Learning (ICML)*, 2018.



Erick Huamán Reyna, Iago Pachêco Gomes, Luis Alberto Rosero, Valdir Grassi, and Denis Fernando Wolf.

A review on multi-scenario decision-making for urban autonomous driving.
IEEE Access, 2026.



David Isele, Reza Rahimi, Akansel Cosgun, Kaushik Subramanian, and Kikuo Fujimura.

Navigating occluded intersections with autonomous vehicles using deep reinforcement learning.
In *IEEE International Conference on Robotics and Automation (ICRA)*, 2018.



Shu Kai et al.

A multi-task reinforcement learning approach for navigating unsignalized intersections.
In *IEEE Intelligent Vehicles Symposium (IV)*, 2020.



Marvin Klimke, Benjamin Gerigk, Benjamin Völz, and Michael Buchholz.

An enhanced graph representation for machine learning based automatic intersection management.
In *IEEE International Intelligent Transportation Systems Conference (ITSC)*, 2022.

References II



Jackeline Rios-Torres and Andreas A. Malikopoulos.

A survey on the coordination of connected and automated vehicles at intersections and merging at highway on-ramps.
IEEE Transactions on Intelligent Transportation Systems, 18(5):1066–1077, 2017.



Maximilian Schier, Christoph Reinders, and Bodo Rosenhahn.

Deep reinforcement learning for autonomous driving using high-level heterogeneous graph representations.
In *IEEE International Conference on Robotics and Automation (ICRA)*, 2023.



Hong Shu, Teng Liu, Xinyu Mu, and Dongpu Cao.

Driving tasks transfer using deep reinforcement learning for decision-making of autonomous vehicles in unsignalized intersection.
IEEE Transactions on Vehicular Technology, 71(1):41–52, 2022.



Martin Treiber, Ansgar Hennecke, and Dirk Helbing.

Congested traffic states in empirical observations and microscopic simulations.
Physical Review E, 62(2):1805, 2000.



Ashish Vaswani et al.

Attention is all you need.
Advances in Neural Information Processing Systems, 2017.



Yiyang Wu, He Chen, and Feng Zhu.

Dcl-aim: Decentralized coordination learning of autonomous intersection management for connected and automated vehicles.
Transportation Research Part C: Emerging Technologies, 103:246–260, 2019.

References III



Wenxuan Xiao, Yuyou Yang, Xinyu Mu, Yi Xie, Xiaolin Tang, Dongpu Cao, and Teng Liu.

Decision-making for autonomous vehicles in random task scenarios at unsignalized intersection using deep reinforcement learning.

IEEE Transactions on Vehicular Technology, 73(6):7812–7825, 2024.



Zijia Zhong, Mark Nejad, and Earl E. Lee.

Autonomous and semiautonomous intersection management: A survey.

IEEE Intelligent Transportation Systems Magazine, 13(2):53–70, 2021.